


SAP AI — Interview Prep Deep-Dive Pack — Index & Glossary


This pack goes deeper than a quick cheat sheet: every PDF here covers **every important file** in one template folder, with working code, a line-by-line explanation of what each part does, and every key point worth remembering for an interview. Read 01 → 10 in order — each stage is a dependency of the next, exactly like the folders in your SAP AI Templates project.

Reading Order & What Each Covers

#	Sheet	Files covered
01	HANA DB Connection	connect_env.py, connect_btp_cf.py, connect_hana_ml.py, health_check.py
02	LLM Setup — Gen AI Hub	llm_setup.py (chat model + embedding model), proxy client internals
03	NL → SQL (HANA AI Query)	schema_inspector.py, sql_generator.py, sql_executor.py, query_agent.py, query_with_summary.py
04	HANA Vector Store	embedding_setup.py, document_loaders.py, text_chunker.py, vector_table.py, embed_and_store.py, similarity_search.py, vector_store_agent.py
05	RAG Pipeline	rag_prompts.py, context_retriever.py, rag_pipeline.py, conversation_memory.py, rag_agent.py
06	SAP Chatbot	chatbot_config.py, prompt_builder.py, rag_retriever.py, sql_retriever.py, chatbot_core.py, chatbot_presets.py, cli_runner.py
07	Orchestration — Core	orchestration_setup.py, templating_module.py, llm_module.py, content_filtering_module.py, data_masking_module.py
08	Orchestration — Advanced	document_grounding_module.py, structured_output_module.py, conversation_history_module.py, orchestration_pipeline.py, orchestration_agent.py
09	CAP + HANA Architecture	db/schema.cds, srv/design-time.cds, CQL query patterns, security model
10	CAP + AI Core + UI5	srv/runtime.js (AI Core calls), UI5 manifest/controller, mta.yaml deployment

The Big Picture — How Every Layer Connects





Glossary — Terms You Must Be Able to Define Cold

Term	One-line definition
SAP AI Core Deployment	BTP service that hosts/serves/orchestrates AI models. A running, callable instance of a model or scenario — has an ID and URL.
Scenario	Category of AI Core workload (foundation-models, orchestration, etc.).
Resource group	Isolation boundary inside one AI Core instance.
Gen AI Hub proxy	Harmonized direct-call API to one foundation model at a time.
Orchestration Service	Managed pipeline: masking → filtering → templating → grounding → LLM → filtering → unmask.
REAL_VECTOR / Vector(n)	HANA's native embedding column type; n = dimension.
COSINE_SIMILARITY()	HANA SQL function scoring 0 (unrelated) to 1 (identical) between two vectors.
CAP	Cloud Application Programming model — CDS + Node.js/Java services framework.
CDS / CQL	Core Data Services (modeling language) / CDS Query Language (its ORM-style query syntax).
HDI container	Isolated HANA schema+artifacts owned by one CAP app.
XSUAA	BTP's OAuth2/JWT service securing CAP services and UI5 apps.
MTA	Multi-Target Application — one deployable bundle of srv+db+app modules.
Destination service	BTP service storing target-system URL+auth so code never hardcodes secrets.

Two Ways to Call an LLM — Comparison You Should Be Able to Explain

	Gen AI Hub Proxy (sheets 02-06)	Orchestration Service (sheets 07-08)
What you get	One model, one API call	A configurable pipeline around the model
Masking / filtering / RAG	You write it yourself	Built into platform config

	Gen AI Hub Proxy (sheets 02-06)	Orchestration Service (sheets 07-08)
State	Stateless — you manage history client-side either way	Same — stateless
Best for	Simple, single-model calls, fast prototyping	Governed, auditable, multi-step pipelines